

**DP-900.examcollection.premium.exam.66q**

Number: DP-900  
Passing Score: 800  
Time Limit: 120 min  
File Version: 1.0

**ExamCollection**

**DP-900**

**Microsoft Azure Data Fundamentals (beta)**

**Version 1.0**

**Sections**

1. Describe core data concepts
2. Describe how to work with relational data on Azure
3. Describe how to work with non-relational data on Azure
4. Describe an analytics workload on Azure

## Exam A

### QUESTION 1 HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

Answer Area

Descriptive analytics tells you

	▼
what is most likely to occur in the future.	
what occurred in the past.	
which actions you can perform to affect outcomes.	
why something occurred in the past.	

Correct Answer:

Answer Area

Descriptive analytics tells you

	▼
what is most likely to occur in the future.	
what occurred in the past.	
which actions you can perform to affect outcomes.	
why something occurred in the past.	

Section: Describe core data concepts  
Explanation

Explanation/Reference:

Reference:

<https://demand-planning.com/2020/01/20/the-differences-between-descriptive-diagnostic-predictive-cognitive-analytics/>

### QUESTION 2 HOTSPOT

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

**NOTE:** Each correct selection is worth one point.

Hot Area:

### Answer Area

Statements	Yes	No
Normalization involves eliminating relationships between database tables.	<input type="radio"/>	<input type="radio"/>
Normalizing a database reduces data redundancy.	<input type="radio"/>	<input type="radio"/>
Normalization improves data integrity.	<input type="radio"/>	<input type="radio"/>

Correct Answer:

### Answer Area

Statements	Yes	No
Normalization involves eliminating relationships between database tables.	<input type="radio"/>	<input checked="" type="radio"/>
Normalizing a database reduces data redundancy.	<input checked="" type="radio"/>	<input type="radio"/>
Normalization improves data integrity.	<input checked="" type="radio"/>	<input type="radio"/>

Section: Describe core data concepts  
Explanation

Explanation/Reference:

Reference:

<https://www.sqlshack.com/what-is-database-normalization-in-sql-server/>

### QUESTION 3

#### HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

### Answer Area

An extract, transform, and load (ETL) process  
requires

- |  |   |
|--|---|
|  | ▼ |
| a matching schema in the data source and the data target.                  |   |
| a target data store powerful enough to transform data.                     |   |
| data that is fully processed before being loaded to the target data store. |   |
| that the data target be a relational database.                             |   |

Correct Answer:

### Answer Area

An extract, transform, and load (ETL) process  
requires

- |  |   |
|--|---|
|  | ▼ |
| a matching schema in the data source and the data target.                  |   |
| a target data store powerful enough to transform data.                     |   |
| data that is fully processed before being loaded to the target data store. |   |
| that the data target be a relational database.                             |   |

Section: Describe core data concepts

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/azure/architecture/data-guide/relational-data/etl>

### QUESTION 4

HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

### Answer Area

In batch processing,

- |  |   |
|--|---|
|  | ▼ |
| data is always inserted one row at a time. |   |
| data is processed in real-time.            |   |
| latency is expected.                       |   |
| processing can only execute serially.      |   |

Correct Answer:

## Answer Area

In batch processing,

	▼
data is always inserted one row at a time.	
data is processed in real-time.	
latency is expected.	
processing can only execute serially.	

### Section: Describe core data concepts

#### Explanation

#### Explanation/Reference:

Reference:

<https://www.bmc.com/blogs/what-is-batch-processing-batch-processing-explained/>

### QUESTION 5

#### DRAG DROP

Match the types of analytics that can be used to answer the business questions.

To answer, drag the appropriate analytics type from the column on the left to its question on the right. Each analytics type may be used once, more than once, or not at all.

**NOTE:** Each correct match is worth one point.

#### Select and Place:

##### Analytics Types

Cognitive
Diagnostic
Descriptive
Predictive
Prescriptive

##### Answer Area

	Why did sales increase last month?
	How do I allocate my budget to buy different inventory items?
	Which people are mentioned in a company's business documents?

#### Correct Answer:

##### Analytics Types

Cognitive
Diagnostic
Descriptive
Predictive
Prescriptive

##### Answer Area

Diagnostic	Why did sales increase last month?
Predictive	How do I allocate my budget to buy different inventory items?
Cognitive	Which people are mentioned in a company's business documents?

### Section: Describe core data concepts

## Explanation

### Explanation/Reference:

Reference:

<https://demand-planning.com/2020/01/20/the-differences-between-descriptive-diagnostic-predictive-cognitive-analytics/>

## QUESTION 6

### HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

### Answer Area

Transcribing audio files is an example of 

	▼
cognitive	
descriptive	
predictive	
prescriptive	

 analytics.

Correct Answer:

### Answer Area

Transcribing audio files is an example of 

	▼
cognitive	
descriptive	
predictive	
prescriptive	

 analytics.

## Section: Describe core data concepts

### Explanation

### Explanation/Reference:

Reference:

<https://azure.microsoft.com/en-us/services/cognitive-services/speech-services/>

## QUESTION 7

### HOTSPOT

You have the following JSON document.

```
"customer" : {
  "first name" : "Ben",
  "last name" : "Smith",
  "address" : {
    "line 1" : "161 Azure Ln",
    "line 2" : "Palo Alto",
    "ZIP code" : "54762"
  },
  "social media": [
    {
      "service" : "twitter",
      "handle" : "@bensmith"
    },
    {
      "service" : "linkedin",
      "handle" : "bensmith"
    }
  ],
  "phone numbers": [
    {
      "type" : "mobile",
      "number" : "555-555-555"
    }
  ]
}
```

Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the JSON document.

**NOTE:** Each correct selection is worth one point.

**Hot Area:**



## Answer Area

Customer is [answer choice].

	▼
a nested array	
a nested object	
a root object	

Address is [answer choice].

	▼
a nested array	
a nested object	
a root object	

Social media is [answer choice].

	▼
a nested array	
a nested object	
a root object	

Correct Answer:

## Answer Area

Customer is [answer choice].

	▼
a nested array	
a nested object	
a root object	

Address is [answer choice].

	▼
a nested array	
a nested object	
a root object	

Social media is [answer choice].

	▼
a nested array	
a nested object	
a root object	

Section: Describe core data concepts  
Explanation



**Explanation/Reference:**

Reference:


[https://www.w3schools.com/js/js\\_json\\_arrays.asp](https://www.w3schools.com/js/js_json_arrays.asp)

[https://www.w3schools.com/js/js\\_json\\_objects.asp](https://www.w3schools.com/js/js_json_objects.asp)

**QUESTION 8**

**HOTSPOT**

You are reviewing the data model shown in the following exhibit.



Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the graphic.

**NOTE:** Each correct selection is worth one point

**Hot Area:**

**Answer Area**

The data model is a [answer choice].

	▼
transactional model	
star schema	
snowflake schema	

Customer is a [answer choice] table.

	▼
fact	
dimension	
bridge	

Correct Answer:

### Answer Area

The data model is a [answer choice].

	▼
transactional model	
star schema	
snowflake schema	

Customer is a [answer choice] table.

	▼
fact	
dimension	
bridge	

Section: Describe core data concepts

Explanation

Explanation/Reference:

Reference:

[https://en.wikipedia.org/wiki/Snowflake\\_schema](https://en.wikipedia.org/wiki/Snowflake_schema)

<https://azure.microsoft.com/en-us/blog/data-models-within-azure-analysis-services-and-power-bi/>

### QUESTION 9

HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

### Answer Area

The massively parallel processing (MPP) engine of Azure Synapse Analytics

	▼
distributes processing across compute nodes.	
distributes processing across control nodes.	
redirects client connections across compute nodes.	
redirects client connections across control nodes.	

Correct Answer:

## Answer Area

The massively parallel processing (MPP) engine of Azure Synapse Analytics

	▼
distributes processing across compute nodes.	
distributes processing across control nodes.	
redirects client connections across compute nodes.	
redirects client connections across control nodes.	

**Section:** Describe core data concepts

**Explanation**

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/azure/synapse-analytics/sql-data-warehouse/massively-parallel-processing-mpp-architecture>

### QUESTION 10

HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

**Hot Area:**

## Answer Area

	▼
A clustered index	
A FileTable	
A foreign key	
A stored procedure	

is an object associated with a table that sorts and stores the data rows in the table based on their key values.

**Correct Answer:**

## Answer Area

	▼
A clustered index	
A FileTable	
A foreign key	
A stored procedure	

is an object associated with a table that sorts and stores the data rows in the table based on their key values.

**Section:** Describe core data concepts

**Explanation**

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/sql/relational-databases/indexes/clustered-and-nonclustered-indexes-described?view=sql-server-ver15>**QUESTION 11**

HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

**Hot Area:****Answer Area**

A relational database must be used when

	▼
a dynamic schema is required.	
data will be stored as key/value pairs.	
storing large images and videos.	
strong consistency guarantees are required.	

**Correct Answer:****Answer Area**

A relational database must be used when

	▼
a dynamic schema is required.	
data will be stored as key/value pairs.	
storing large images and videos.	
strong consistency guarantees are required.	

**Section: Describe how to work with relational data on Azure****Explanation****Explanation/Reference:****QUESTION 12**

HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

**Hot Area:****Answer Area**

Relational data uses

	▼
collections	
columns	
keys	
partitions	

to enforce relationships between different tables.

**Correct Answer:**

### Answer Area

Relational data uses 

▼
collections
columns
keys
partitions

 to enforce relationships between different tables.

▼
collections
columns
keys
partitions

### Section: Describe how to work with relational data on Azure

#### Explanation

#### Explanation/Reference:

Reference:

<https://teachcomputerscience.com/relational-databases/>

### QUESTION 13

#### HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

#### Hot Area:

### Answer Area

▼	is a virtual table that contains content defined by a query.
A heap	
A stored procedure	
A view	
An index	

#### Correct Answer:

### Answer Area

▼	is a virtual table that contains content defined by a query.
A heap	
A stored procedure	
A view	
An index	

### Section: Describe how to work with relational data on Azure

#### Explanation

#### Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/sql/relational-databases/views/views>

### QUESTION 14

You need to query a table named Products in an Azure SQL database.

Which three requirements must be met to query the table from the internet? Each correct answer presents part

of the solution. (Choose three.)

**NOTE:** Each correct selection is worth one point.

- A. You must be assigned the Reader role for the resource group that contains the database.
- B. You must have `SELECT` access to the Products table.
- C. You must have a user in the database.
- D. You must be assigned the Contributor role for the resource group that contains the database.
- E. Your IP address must be allowed to connect to the database.

**Correct Answer:** ACE

**Section:** Describe how to work with relational data on Azure

**Explanation**

**Explanation/Reference:**

#### QUESTION 15

You have an inventory management database that contains the following table.

ProductName	Quantity
Product1	100
Product2	129
Product3	176

Which statement should you use in a SQL query to change the inventory quantity of Product1 to 270?

- A. `INSERT`
- B. `MERGE`
- C. `UPDATE`
- D. `CREATE`

**Correct Answer:** C

**Section:** Describe how to work with relational data on Azure

**Explanation**

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/sql/t-sql/queries/update-transact-sql?view=sql-server-ver15>

#### QUESTION 16

Your company needs to implement a relational database in Azure. The solution must minimize ongoing maintenance.

Which Azure service should you use?

- A. Azure HDInsight
- B. Azure SQL Database
- C. Azure Cosmos DB
- D. SQL Server on Azure virtual machines

**Correct Answer:** B

**Section: Describe how to work with relational data on Azure**

**Explanation**

**Explanation/Reference:**

Reference:

<https://azure.microsoft.com/en-us/services/sql-database/#features>

**QUESTION 17**

You are writing a set of SQL queries that administrators will use to troubleshoot an Azure SQL database.

You need to embed documents and query results into a SQL notebook.

What should you use?

- A. Microsoft SQL Server Management Studio (SSMS)
- B. Azure Data Studio
- C. Azure CLI
- D. Azure PowerShell

**Correct Answer: B**

**Section: Describe how to work with relational data on Azure**

**Explanation**

**Explanation/Reference:**

Reference:

<https://www.mssqltips.com/sqlservertip/5997/create-sql-server-notebooks-in-azure-data-studio/>

**QUESTION 18**

**DRAG DROP**

Match the terms to the appropriate descriptions.

To answer, drag the appropriate term from the column on the left to its description on the right. Each term may be used once, more than once, or not at all.

**NOTE:** Each correct match is worth one point.

**Select and Place:**

**Terms**

Index
View
Table

**Answer Area**

	A database object that holds data
	A database object whose content is defined by a query
	A database object that helps improve the speed of data retrieval

**Correct Answer:**

**Terms**

Index
View
Table

**Answer Area**

Table	A database object that holds data
View	A database object whose content is defined by a query
Index	A database object that helps improve the speed of data retrieval

**Section: Describe how to work with relational data on Azure**



## Explanation

### Explanation/Reference:

Reference:

[https://en.wikipedia.org/wiki/Table\\_\(database\)#:~:text=A%20table%20is%20a%20collection,table%20format%20within%20a%20database.&text=In%20relational%20databases%2C%20and%20flat,a%20row%20and%20column%20intersect.](https://en.wikipedia.org/wiki/Table_(database)#:~:text=A%20table%20is%20a%20collection,table%20format%20within%20a%20database.&text=In%20relational%20databases%2C%20and%20flat,a%20row%20and%20column%20intersect.)

[https://en.wikipedia.org/wiki/View\\_\(SQL\)](https://en.wikipedia.org/wiki/View_(SQL))

[https://en.wikipedia.org/wiki/Database\\_index#:~:text=A%20database%20index%20is%20a,maintain%20the%20index%20data%20structure.](https://en.wikipedia.org/wiki/Database_index#:~:text=A%20database%20index%20is%20a,maintain%20the%20index%20data%20structure.)

### QUESTION 19

Which command-line tool can you use to query Azure SQL databases?

- A. sqlcmd
- B. bcp
- C. azdata
- D. Azure CLI

**Correct Answer: D**

**Section: Describe how to work with relational data on Azure**

### Explanation

### Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/azure/azure-sql/database/az-cli-script-samples-content-guide?tabs=single-database>

### QUESTION 20

#### HOTSPOT

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

**NOTE:** Each correct selection is worth one point.

**Hot Area:**

#### Answer Area

Statements	Yes	No
Azure SQL Database includes a fully managed backup service.	<input type="radio"/>	<input type="radio"/>
Azure SQL Database has built-in high availability.	<input type="radio"/>	<input type="radio"/>
Azure SQL Database can use Azure Advanced Threat Protection (ATP).	<input type="radio"/>	<input type="radio"/>

**Correct Answer:**

### Answer Area

Statements	Yes	No
Azure SQL Database includes a fully managed backup service.	<input checked="" type="radio"/>	<input type="radio"/>
Azure SQL Database has built-in high availability.	<input checked="" type="radio"/>	<input type="radio"/>
Azure SQL Database can use Azure Advanced Threat Protection (ATP).	<input checked="" type="radio"/>	<input type="radio"/>

**Section: Describe how to work with relational data on Azure**

**Explanation**

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/azure/azure-sql/database/sql-database-paas-overview>

<https://azure.microsoft.com/en-us/blog/announcing-sql-atp-and-sql-vulnerability-assessment-general-availability/>

### QUESTION 21

HOTSPOT

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

**NOTE:** Each correct selection is worth one point.

**Hot Area:**

### Answer Area

Statements	Yes	No
You can use Azure Data Studio to query a Microsoft SQL Server big data cluster.	<input type="radio"/>	<input type="radio"/>
You can use Microsoft SQL Server Management Studio (SSMS) to query an Azure Synapse Analytics data warehouse.	<input type="radio"/>	<input type="radio"/>
You can use MySQL Workbench to query Azure Database for MariaDB databases.	<input type="radio"/>	<input type="radio"/>

**Correct Answer:**

## Answer Area

Statements	Yes	No
You can use Azure Data Studio to query a Microsoft SQL Server big data cluster.	<input checked="" type="radio"/>	<input type="radio"/>
You can use Microsoft SQL Server Management Studio (SSMS) to query an Azure Synapse Analytics data warehouse.	<input checked="" type="radio"/>	<input type="radio"/>
You can use MySQL Workbench to query Azure Database for MariaDB databases.	<input checked="" type="radio"/>	<input type="radio"/>

**Section: Describe how to work with relational data on Azure**  
**Explanation**

### Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/sql/big-data-cluster/connect-to-big-data-cluster?view=sql-server-ver15>

<https://docs.microsoft.com/en-us/azure/synapse-analytics/sql-data-warehouse/sql-data-warehouse-query-ssms>

<https://docs.microsoft.com/en-us/azure/mariadb/connect-workbench>

### QUESTION 22

#### HOTSPOT

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

**NOTE:** Each correct selection is worth one point.

**Hot Area:**

## Answer Area

Statements	Yes	No
Platform as a service (PaaS) database offerings in Azure provide built-in high availability.	<input type="radio"/>	<input type="radio"/>
Platform as a service (PaaS) database offerings in Azure provide configurable scaling options.	<input type="radio"/>	<input type="radio"/>
Platform as a service (PaaS) database offerings in Azure reduce the administrative overhead for managing hardware.	<input type="radio"/>	<input type="radio"/>

**Correct Answer:**

## Answer Area

Statements	Yes	No
Platform as a service (PaaS) database offerings in Azure provide built-in high availability.	<input checked="" type="radio"/>	<input type="radio"/>
Platform as a service (PaaS) database offerings in Azure provide configurable scaling options.	<input checked="" type="radio"/>	<input type="radio"/>
Platform as a service (PaaS) database offerings in Azure reduce the administrative overhead for managing hardware.	<input checked="" type="radio"/>	<input type="radio"/>

### Section: Describe how to work with relational data on Azure

#### Explanation

#### Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/azure/azure-sql/database/sql-database-paas-overview>

### QUESTION 23

#### HOTSPOT

You have the following SQL query.

```
INSERT INTO dbo.Products (ProductID, ProductName, Price, ProductDescription)
VALUES (1, 'Clamp', 12.48, 'Workbench clamp') ;
```

What are `dbo.Products` and `ProductName`? To answer, select the appropriate options in the answer area.

**NOTE:** Each correct selection is worth one point.

#### Hot Area:

## Answer Area

dbo.Products : 

	▼
A column	
A database	
A table	
An index	

ProductName : 

	▼
A column	
A database	
A table	
An index	

Correct Answer:

## Answer Area

dbo.Products : 

	▼
A column	
A database	
A table	
An index	

ProductName : 

	▼
A column	
A database	
A table	
An index	

Section: Describe how to work with relational data on Azure

Explanation

Explanation/Reference:

### QUESTION 24

HOTSPOT

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

**NOTE:** Each correct selection is worth one point.

Hot Area:

**Answer Area**

Statements	Yes	No
You must apply patches to Azure SQL databases regularly.	<input type="radio"/>	<input type="radio"/>
You need a Microsoft 365 subscription to create an Azure SQL database.	<input type="radio"/>	<input type="radio"/>
You can use existing Microsoft SQL Server licenses to reduce the cost of Azure SQL databases.	<input type="radio"/>	<input type="radio"/>

Correct Answer:

**Answer Area**

Statements	Yes	No
You must apply patches to Azure SQL databases regularly.	<input type="radio"/>	<input checked="" type="radio"/>
You need a Microsoft 365 subscription to create an Azure SQL database.	<input type="radio"/>	<input checked="" type="radio"/>
You can use existing Microsoft SQL Server licenses to reduce the cost of Azure SQL databases.	<input checked="" type="radio"/>	<input type="radio"/>

**Section: Describe how to work with relational data on Azure**

**Explanation**

**Explanation/Reference:**

Reference:

<https://azure.microsoft.com/en-gb/blog/hot-patching-sql-server-engine-in-azure-sql-database/>

<https://azure.microsoft.com/en-us/services/sql-database/#product-overview>

**QUESTION 25**

Which statement is an example of Data Definition Language (DDL)?

- A. SELECT
- B. JOIN
- C. MERGE
- D. CREATE

**Correct Answer: D**

**Section: Describe how to work with relational data on Azure**

**Explanation**

**Explanation/Reference:**

Reference:

<https://www.geeksforgeeks.org/difference-between-ddl-and-dml-in-dbms/>

### QUESTION 26

#### HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

**Hot Area:**

Answer Area

Transparent Data Encryption (TDE) encrypts

	▼
a column to protect data at rest and in transit.	
queries and their results in order to protect data in transit.	
the database to protect data at rest.	
the server to protect data at rest.	

**Correct Answer:**

Answer Area

Transparent Data Encryption (TDE) encrypts

	▼
a column to protect data at rest and in transit.	
queries and their results in order to protect data in transit.	
the database to protect data at rest.	
the server to protect data at rest.	

### Section: Describe how to work with relational data on Azure

#### Explanation

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/azure/azure-sql/database/transparent-data-encryption-tde-overview?tabs=azure-portal>

### QUESTION 27

You need to ensure that users use multi-factor authentication (MFA) when connecting to an Azure SQL database.

Which type of authentication should you use?

- A. service principal authentication
- B. Azure Active Directory (Azure AD) authentication
- C. SQL authentication
- D. certificate authentication

**Correct Answer: B**

### Section: Describe how to work with relational data on Azure

#### Explanation

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/azure/azure-sql/database/authentication-mfa-ssms-overview>

### QUESTION 28

What is a benefit of hosting a database on Azure SQL managed instance as compared to an Azure SQL



database?

- A. built-in high availability
- B. native support for cross-database queries and transactions
- C. system-initiated automatic backups
- D. support for encryption at rest

**Correct Answer:** B

**Section:** Describe how to work with relational data on Azure

**Explanation**

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/azure/azure-sql/database/features-comparison>

### QUESTION 29

#### HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

**Hot Area:**

**Answer Area**

By default, each Azure SQL database is protected by

	▼
a network security group (NSG).	
a server-level firewall.	
Azure Firewall.	
Azure Front Door.	

**Correct Answer:**

**Answer Area**

By default, each Azure SQL database is protected by

	▼
a network security group (NSG).	
a server-level firewall.	
Azure Firewall.	
Azure Front Door.	

**Section:** Describe how to work with relational data on Azure

**Explanation**

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/azure/security/fundamentals/infrastructure-sql>

### QUESTION 30

You need to design and model a database by using a graphical tool that supports project-oriented offline database development.

What should you use?

- A. Microsoft SQL Server Data Tools (SSDT)

- B. Microsoft SQL Server Management Studio (SSMS)
- C. Azure Databricks
- D. Azure Data Studio

**Correct Answer:** A

**Section:** Describe how to work with relational data on Azure

**Explanation**

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/sql/ssdt/project-oriented-offline-database-development?view=sql-server-ver15>

### QUESTION 31

DRAG DROP

Match the security components to the appropriate scenarios.

To answer, drag the appropriate component from the column on the left to its scenario on the right. Each component may be used once, more than once, or not at all.

**NOTE:** Each correct match is worth one point.

**Select and Place:**

**Components**

Authentication

Firewall

Encryption

**Answer Area**

Prevent access to an Azure SQL database from another network.

Support Azure Active Directory (Azure AD) sign-ins to an Azure SQL database.

Ensure that sensitive data never appears as plain text in an Azure SQL database.

**Correct Answer:**

**Components**

Authentication

Firewall

Encryption

**Answer Area**

Firewall

Prevent access to an Azure SQL database from another network.

Authentication

Support Azure Active Directory (Azure AD) sign-ins to an Azure SQL database.

Encryption

Ensure that sensitive data never appears as plain text in an Azure SQL database.

**Section:** Describe how to work with relational data on Azure

**Explanation**

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/sql/database-engine/configure-windows/configure-a-windows-firewall-for-database-engine-access?view=sql-server-ver15>

<https://docs.microsoft.com/en-us/azure/azure-sql/database/authentication-aad-overview>

<https://docs.microsoft.com/en-us/azure/azure-sql/database/always-encrypted-certificate-store-configure>

### QUESTION 32

You have a transactional application that stores data in an Azure SQL managed instance.

When should you implement a read-only database replica?

- A. You need to generate reports without affecting the transactional workload.

- B. You need to audit the transactional application.
- C. You need to implement high availability in the event of a regional outage.
- D. You need to improve the recovery point objective (RPO).

**Correct Answer:** C

**Section:** Describe how to work with relational data on Azure

**Explanation**

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/sql/database-engine/availability-groups/windows/active-secondaries-readable-secondary-replicas-always-on-availability-groups?view=sql-server-ver15>

### QUESTION 33

At which two levels can you set the throughput for an Azure Cosmos DB account? Each correct answer presents a complete solution. (Choose two.)

**NOTE:** Each correct selection is worth one point.

- A. database
- B. item
- C. container
- D. partition

**Correct Answer:** AC

**Section:** Describe how to work with non-relational data on Azure

**Explanation**

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/azure/cosmos-db/set-throughput>

### QUESTION 34

HOTSPOT

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

**NOTE:** Each correct selection is worth one point.

**Hot Area:**

**Answer Area**

Statements	Yes	No
Azure Table storage supports multiple read replicas.	<input type="radio"/>	<input type="radio"/>
Azure Table storage supports multiple write regions.	<input type="radio"/>	<input type="radio"/>
The Azure Cosmos DB Table API supports multiple read replicas.	<input type="radio"/>	<input type="radio"/>
The Azure Cosmos DB Table API supports multiple write regions.	<input type="radio"/>	<input type="radio"/>

**Correct Answer:**

## Answer Area

Statements	Yes	No
Azure Table storage supports multiple read replicas.	<input checked="" type="radio"/>	<input type="radio"/>
Azure Table storage supports multiple write regions.	<input type="radio"/>	<input checked="" type="radio"/>
The Azure Cosmos DB Table API supports multiple read replicas.	<input checked="" type="radio"/>	<input type="radio"/>
The Azure Cosmos DB Table API supports multiple write regions.	<input checked="" type="radio"/>	<input type="radio"/>

### Section: Describe how to work with non-relational data on Azure Explanation

#### Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/azure/cosmos-db/high-availability>

### QUESTION 35

#### DRAG DROP

Match the types of data stores to the appropriate scenarios.

To answer, drag the appropriate data store type from the column on the left to its scenario on the right. Each data store type may be used once, more than once, or not at all.

**NOTE:** Each correct match is worth one point.

#### Select and Place:

##### Data Store Types

Graph
Object
Key/value

##### Answer Area

	Application users and their default language
	Medical images and their associated metadata
	Employee data that shows the relationships between employees

#### Correct Answer:

##### Data Store Types

Graph
Object
Key/value

##### Answer Area

Key/value	Application users and their default language
Object	Medical images and their associated metadata
Graph	Employee data that shows the relationships between employees

### Section: Describe how to work with non-relational data on Azure Explanation

#### Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/azure/architecture/guide/technology-choices/data-store-overview>

**QUESTION 36**

You have an Azure Cosmos DB account that uses the Core (SQL) API.

Which two settings can you configure at the container level? Each correct answer presents a complete solution. (Choose two.)

**NOTE:** Each correct selection is worth one point.

- A. the throughput
- B. the read region
- C. the partition key
- D. the API

**Correct Answer:** AC

**Section:** Describe how to work with non-relational data on Azure

**Explanation**

**Explanation/Reference:**

Reference:

<https://www.sqlshack.com/start-your-journey-with-azure-cosmos-db/>

**QUESTION 37**

Your company is designing a data store that will contain student data. The data has the following format.

StudentNumber	StudentInformation
7634634	First name: Ben Last: Smith Preferred Name: Benjamin
7634634	First Name: Dominik Last Name: Paiha Email Address: dpaiha@contoso.com MCP ID: 931817
7634636	First Name: Reshma Last Name: Patel Phone number: 514-555-1101
7634637	First Name: Yun-Feng Last Name: Peng

Which type of data store should you use?

- A. graph
- B. key/value
- C. object
- D. columnar

**Correct Answer:** D

**Section:** Describe how to work with non-relational data on Azure

**Explanation**

**Explanation/Reference:**

#### **QUESTION 38**

Which storage solution supports role-based access control (RBAC) at the file and folder level?

- A. Azure Disk Storage
- B. Azure Data Lake Storage
- C. Azure Blob storage
- D. Azure Queue storage

**Correct Answer:** B

**Section:** Describe how to work with non-relational data on Azure

**Explanation**

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/azure/storage/blobs/data-lake-storage-access-control>

#### **QUESTION 39**

You need to store data in Azure Blob storage for seven years to meet your company's compliance requirements. The retrieval time of the data is unimportant. The solution must minimize storage costs.

Which storage tier should you use?

- A. Archive
- B. Hot
- C. Cool

**Correct Answer:** A

**Section:** Describe how to work with non-relational data on Azure

**Explanation**

**Explanation/Reference:**

Reference:

[https://cloud.netapp.com/blog/azure-blob-storage-pricing-the-complete-guide-azure-cvo-blg#H1\\_4](https://cloud.netapp.com/blog/azure-blob-storage-pricing-the-complete-guide-azure-cvo-blg#H1_4)

#### **QUESTION 40**

Which type of non-relational data store supports a flexible schema, stores data as JSON files, and stores the all the data for an entity in the same document?

- A. document
- B. columnar
- C. graph
- D. time series

**Correct Answer:** A

**Section:** Describe how to work with non-relational data on Azure

**Explanation**

**Explanation/Reference:**

#### QUESTION 41

##### DRAG DROP

Match the Azure Cosmos DB APIs to the appropriate data structures.

To answer, drag the appropriate API from the column on the left to its data structure on the right. Each API may be used once, more than once, or not at all.

**NOTE:** Each correct match is worth one point.

Select and Place:

APIs	Answer Area
Cassandra API	<input type="text"/> Graph data
Gremlin API	<input type="text"/> JSON documents
MongoDB API	<input type="text"/> Key/value data
Table API	

Correct Answer:

APIs	Answer Area
Cassandra API	<input type="text"/>
Gremlin API	Gremlin API Graph data
Gremlin API	MongoDB API JSON documents
MongoDB API	Table API Key/value data
Table API	

**Section:** Describe how to work with non-relational data on Azure

**Explanation**

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/azure/cosmos-db/faq>

#### QUESTION 42

##### HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

**Hot Area:**



### Answer Area

To configure an Azure Storage account to support both security at the folder level and atomic directory manipulation,

- enable the hierarchical namespace.
- set Account kind to BlobStorage.
- set Performance to Premium.
- set Replication to Read-access geo-redundant storage (RA-GRS).

**Correct Answer:**

### Answer Area

To configure an Azure Storage account to support both security at the folder level and atomic directory manipulation,

- enable the hierarchical namespace.
- set Account kind to BlobStorage.
- set Performance to Premium.
- set Replication to Read-access geo-redundant storage (RA-GRS).

**Section: Describe how to work with non-relational data on Azure**

### Explanation

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/azure/storage/blobs/data-lake-storage-namespace>

### QUESTION 43

#### HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

**Hot Area:**

### Answer Area

You can query a graph database in Azure Cosmos DB

- as a JSON document by using a SQL-like language.
- as a partitioned row store by using Cassandra Query Language (CQL).
- as a partitioned row store by using Language-Integrated Query (LINQ).
- as nodes and edges by using the Gremlin language.

**Correct Answer:**

### Answer Area

You can query a graph database in Azure Cosmos DB

- as a JSON document by using a SQL-like language.
- as a partitioned row store by using Cassandra Query Language (CQL).
- as a partitioned row store by using Language-Integrated Query (LINQ).
- as nodes and edges by using the Gremlin language.

**Section: Describe how to work with non-relational data on Azure**

### Explanation

**Explanation/Reference:**

Reference:

<https://www.sqlshack.com/graph-database-implementation-with-azure-cosmos-db-using-the-api/>

#### QUESTION 44

You manage an application that stores data in a shared folder on a Windows server.

You need to move the shared folder to Azure Storage.

Which type of Azure Storage should you use?

- A. queue
- B. blob
- C. file
- D. table

**Correct Answer: B**

**Section: Describe how to work with non-relational data on Azure**

**Explanation**

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/azure/storage/common/storage-use-azcopy-files>

#### QUESTION 45

Your company is designing a database that will contain session data for a website. The data will include notifications, personalization attributes, and products that are added to a shopping cart.

Which type of data store will provide the lowest latency to retrieve the data?

- A. key/value
- B. graph
- C. columnar
- D. document

**Correct Answer: C**

**Section: Describe how to work with non-relational data on Azure**

**Explanation**

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/azure/architecture/data-guide/technology-choices/analytical-data-stores>

#### QUESTION 46

HOTSPOT

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

**NOTE:** Each correct selection is worth one point.

**Hot Area:**

### Answer Area

Statements	Yes	No
When ingesting data from Azure Data Lake Storage across Azure regions, you will incur costs for bandwidth.	<input type="radio"/>	<input type="radio"/>
You can use blob, table, and file storage in the same Azure Storage account.	<input type="radio"/>	<input type="radio"/>
You implement Azure Data Lake Storage by creating an Azure Storage account.	<input type="radio"/>	<input type="radio"/>

Correct Answer:

### Answer Area

Statements	Yes	No
When ingesting data from Azure Data Lake Storage across Azure regions, you will incur costs for bandwidth.	<input checked="" type="radio"/>	<input type="radio"/>
You can use blob, table, and file storage in the same Azure Storage account.	<input checked="" type="radio"/>	<input type="radio"/>
You implement Azure Data Lake Storage by creating an Azure Storage account.	<input checked="" type="radio"/>	<input type="radio"/>

**Section: Describe how to work with non-relational data on Azure**

**Explanation**

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/azure/data-lake-store/data-lake-store-get-started-portal>

<https://docs.microsoft.com/en-us/azure/storage/common/storage-account-overview>

<https://azure.microsoft.com/en-us/pricing/details/bandwidth/>

### QUESTION 47

HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

**Hot Area:**

## Answer Area

When using the Azure Cosmos DB Gremlin API, the container resource type is projected as a

	▼
graph.	
table.	
partition key.	
document.	

Correct Answer:

## Answer Area

When using the Azure Cosmos DB Gremlin API, the container resource type is projected as a

	▼
graph.	
table.	
partition key.	
document.	

Section: Describe how to work with non-relational data on Azure

Explanation

Explanation/Reference:

Reference:

<https://docs.microsoft.com/en-us/azure/cosmos-db/create-graph-gremlin-console>

### QUESTION 48

Which scenario is an example of a streaming workload?

- A. sending transactions that are older than a month to an archive
- B. sending transactions daily from point of sale (POS) devices
- C. sending telemetry data from edge devices
- D. sending cloud infrastructure metadata every 30 minutes

Correct Answer: C

Section: Describe an analytics workload on Azure

Explanation

Explanation/Reference:

### QUESTION 49

HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

Hot Area:

Answer Area

Batch workloads

	▼
process data in memory, row-by-row.	
collect and process data at most once a day.	
process data as new data is received in near real-time.	
collect data and then process the data when a condition is met.	

Correct Answer:

Answer Area

Batch workloads

	▼
process data in memory, row-by-row.	
collect and process data at most once a day.	
process data as new data is received in near real-time.	
collect data and then process the data when a condition is met.	

Section: Describe an analytics workload on Azure

Explanation

Explanation/Reference:

#### QUESTION 50

##### HOTSPOT

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

**NOTE:** Each correct selection is worth one point.

Hot Area:

Answer Area

Statements	Yes	No
Processing salary payments once a month is an example of a batch workload.	<input type="radio"/>	<input type="radio"/>
A wind turbine that sends 50 sensor readings per second is an example of a streaming workload.	<input type="radio"/>	<input type="radio"/>
A home electricity meter that sends readings once a day to an energy provider is an example of a streaming workload.	<input type="radio"/>	<input type="radio"/>

Correct Answer:

## Answer Area

Statements	Yes	No
Processing salary payments once a month is an example of a batch workload.	<input checked="" type="radio"/>	<input type="radio"/>
A wind turbine that sends 50 sensor readings per second is an example of a streaming workload.	<input checked="" type="radio"/>	<input type="radio"/>
A home electricity meter that sends readings once a day to an energy provider is an example of a streaming workload.	<input type="radio"/>	<input checked="" type="radio"/>

**Section: Describe an analytics workload on Azure**  
**Explanation**

**Explanation/Reference:**

### QUESTION 51

You need to gather real-time telemetry data from a mobile application.

Which type of workload describes this scenario?

- A. Online Transaction Processing (OLTP)
- B. batch
- C. massively parallel processing (MPP)
- D. streaming

**Correct Answer: D**

**Section: Describe an analytics workload on Azure**  
**Explanation**

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-in/azure/azure-monitor/overview>

### QUESTION 52

You have a SQL pool in Azure Synapse Analytics that is only used actively every night for eight hours.

You need to minimize the cost of the SQL pool during idle times. The solution must ensure that the data remains intact.

What should you do on the SQL pool?

- A. Scale down the data warehouse units (DWUs).
- B. Pause the pool.
- C. Create a user-defined restore point.
- D. Delete the pool

**Correct Answer: B**

**Section: Describe an analytics workload on Azure**  
**Explanation**

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/azure/synapse-analytics/sql-data-warehouse/sql-data-warehouse-manage-compute-overview>

**QUESTION 53**

Which Azure Data Factory component initiates the execution of a pipeline?

- A. a control flow
- B. a trigger
- C. a parameter
- D. an activity

**Correct Answer:** B

**Section:** Describe an analytics workload on Azure

**Explanation**

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/azure/data-factory/concepts-pipeline-execution-triggers#trigger-execution>

**QUESTION 54**

Your company has a reporting solution that has paginated reports. The reports query a dimensional model in a data warehouse.

Which type of processing does the reporting solution use?

- A. stream processing
- B. batch processing
- C. Online Analytical Processing (OLAP)
- D. Online Transaction Processing (OLTP)

**Correct Answer:** C

**Section:** Describe an analytics workload on Azure

**Explanation**

**Explanation/Reference:**

Reference:

<https://datawarehouseinfo.com/how-does-oltp-differ-from-olap-database/>

**QUESTION 55**

**DRAG DROP**

Match the types of activities to the appropriate Azure Data Factory activities.

To answer, drag the appropriate activity type from the column on the left to its Data Factory activity on the right. Each activity type may be used once, more than once, or not at all.

**NOTE:** Each correct match is worth one point.

**Select and Place:**



### Activity Types

Control
Data movement
Data transformation

### Answer Area

	Copy
	Mapping data flow
	Until

**Correct Answer:**

### Activity Types

Control
Data movement
Data transformation

### Answer Area

Data movement	Copy
Data transformation	Mapping data flow
Control	Until

**Section: Describe an analytics workload on Azure**  
**Explanation**

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/azure/data-factory/concepts-pipelines-activities#data-transformation-activities>

#### QUESTION 56

What are three characteristics of an Online Transaction Processing (OLTP) workload? Each correct answer presents a complete solution. (Choose three.)

**NOTE:** Each correct selection is worth one point.

- A. denormalized data
- B. heavy writes and moderate reads
- C. light writes and heavy reads
- D. schema on write
- E. schema on read
- F. normalized data

**Correct Answer:** BDE

**Section: Describe an analytics workload on Azure**  
**Explanation**

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/azure/architecture/data-guide/relational-data/online-transaction-processing>

#### QUESTION 57

Which two activities can be performed entirely by using the Microsoft Power BI service? Each correct answer presents a complete solution. (Choose two.)

**NOTE:** Each correct selection is worth one point.

- A. report and dashboard creation
- B. report sharing and distribution
- C. data modeling
- D. data acquisition and preparation

**Correct Answer:** AD

**Section:** Describe an analytics workload on Azure

**Explanation**

**Explanation/Reference:**

#### QUESTION 58

##### HOTSPOT

To complete the sentence, select the appropriate option in the answer area.

**Hot Area:**

**Answer Area**

In Azure Data Factory, you can use 

	▼
--	---

 to orchestrate pipeline

- |                        |
|------------------------|
| a control flow         |
| a dataset              |
| a linked service       |
| an integration runtime |

activities that depend on the output of other pipeline activities.

**Correct Answer:**

**Answer Area**

In Azure Data Factory, you can use 

	▼
--	---

 to orchestrate pipeline

- |                        |
|------------------------|
| a control flow         |
| a dataset              |
| a linked service       |
| an integration runtime |

activities that depend on the output of other pipeline activities.

**Section:** Describe an analytics workload on Azure

**Explanation**

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/azure/data-factory/frequently-asked-questions>

#### QUESTION 59

You need to develop a solution to provide data to executives. The solution must provide an interactive graphical interface, depict various key performance indicators, and support data exploration by using drill down.

What should you use in Microsoft Power BI?

- A. a dashboard
- B. a report
- C. a dataflow
- D. Microsoft Power Apps

**Correct Answer:** B

**Section:** Describe an analytics workload on Azure

**Explanation**

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/power-bi/consumer/end-user-dashboards>

<https://docs.microsoft.com/en-us/power-bi/visuals/power-bi-visualization-kpi>

<https://docs.microsoft.com/en-us/power-bi/consumer/end-user-drill>

#### QUESTION 60

Which two Azure services can be used to provision Apache Spark clusters? Each correct answer presents a complete solution. (Choose two.)

**NOTE:** Each correct selection is worth one point.

- A. Azure Time Series Insights
- B. Azure HDInsight
- C. Azure Databricks
- D. Azure Log Analytics

**Correct Answer:** BC

**Section:** Describe an analytics workload on Azure

**Explanation**

**Explanation/Reference:**

Reference:

<https://www.sqlshack.com/a-beginners-guide-to-azure-databricks/>

#### QUESTION 61

You have a quality assurance application that reads data from a data warehouse.

Which type of processing does the application use?

- A. Online Transaction Processing (OLTP)
- B. batch processing
- C. Online Analytical Processing (OLAP)
- D. stream processing

**Correct Answer:** A

**Section:** Describe an analytics workload on Azure

**Explanation**

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/azure/architecture/data-guide/relational-data/online-transaction-processing>

**QUESTION 62**

Which three objects can be added to a Microsoft Power BI dashboard? Each correct answer presents a complete solution. (Choose three.)

**NOTE:** Each correct selection is worth one point.

- A. a report page
- B. a Microsoft PowerPoint slide
- C. a visualization from a report
- D. a dataflow
- E. a text box

**Correct Answer:** ACE

**Section:** Describe an analytics workload on Azure

**Explanation**

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/power-bi/consumer/end-user-dashboards>

<https://docs.microsoft.com/en-us/power-bi/create-reports/service-dashboard-add-widget>

**QUESTION 63**

HOTSPOT

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

**NOTE:** Each correct selection is worth one point.

**Hot Area:**

**Answer Area**

Statements	Yes	No
A Microsoft Power BI dashboard is associated with a single workspace.	<input type="radio"/>	<input type="radio"/>
A Microsoft Power BI dashboard can only display visualizations from a single dataset.	<input type="radio"/>	<input type="radio"/>
A Microsoft Power BI dashboard can display visualizations from a Microsoft Excel workbook.	<input type="radio"/>	<input type="radio"/>

**Correct Answer:**

## Answer Area

Statements	Yes	No
A Microsoft Power BI dashboard is associated with a single workspace.	<input type="radio"/>	<input checked="" type="radio"/>
A Microsoft Power BI dashboard can only display visualizations from a single dataset.	<input type="radio"/>	<input checked="" type="radio"/>
A Microsoft Power BI dashboard can display visualizations from a Microsoft Excel workbook.	<input checked="" type="radio"/>	<input type="radio"/>

**Section:** Describe an analytics workload on Azure  
**Explanation**

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/power-bi/connect-data/service-datasets-across-workspaces>

<https://docs.microsoft.com/en-us/power-bi/consumer/end-user-dashboards><https://powerbi.microsoft.com/en-us/excel-and-power-bi/>

### QUESTION 64

Which Azure Data Factory component provides the compute environment for activities?

- A. a linked service
- B. an integration runtime
- C. a control flow
- D. a pipeline

**Correct Answer:** A

**Section:** Describe an analytics workload on Azure  
**Explanation**

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/azure/data-factory/v1/data-factory-compute-linked-services>

### QUESTION 65

You need to use Transact-SQL to query files in Azure Data Lake Storage from an Azure Synapse Analytics data warehouse.

What should you use to query the files?

- A. Azure Functions
- B. Microsoft SQL Server Integration Services (SSIS)
- C. PolyBase
- D. Azure Data Factory

**Correct Answer:** C

**Section: Describe an analytics workload on Azure**  
**Explanation**

**Explanation/Reference:**

Reference:

<https://docs.databricks.com/data/data-sources/azure/synapse-analytics.html>

**QUESTION 66**

**HOTSPOT**

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

**NOTE:** Each correct selection is worth one point.

**Hot Area:**

**Answer Area**

Statements	Yes	No
Azure Databricks is an Apache Spark-based collaborative analytics platform.	<input type="radio"/>	<input type="radio"/>
Azure Analysis Services is used for transactional workloads.	<input type="radio"/>	<input type="radio"/>
Azure Data Factory orchestrates data integration workflows.	<input type="radio"/>	<input type="radio"/>

**Correct Answer:**

**Answer Area**

Statements	Yes	No
Azure Databricks is an Apache Spark-based collaborative analytics platform.	<input checked="" type="radio"/>	<input type="radio"/>
Azure Analysis Services is used for transactional workloads.	<input type="radio"/>	<input checked="" type="radio"/>
Azure Data Factory orchestrates data integration workflows.	<input checked="" type="radio"/>	<input type="radio"/>

**Section: Describe an analytics workload on Azure**  
**Explanation**

**Explanation/Reference:**

Reference:

<https://docs.microsoft.com/en-us/azure/databricks/scenarios/what-is-azure-databricks>

<https://docs.microsoft.com/en-us/azure/analysis-services/analysis-services-overview>

<https://docs.microsoft.com/en-us/azure/data-factory/introduction>